

## CUSTOMER\_CHURN

```

import pandas as pd
import numpy as np

# Load
df = pd.read_csv("customer_level.csv")

# Strip whitespace from string columns
for col in df.select_dtypes(include="object").columns:
    df[col] = df[col].astype(str).str.strip().replace({"nan": np.nan})

# Normalize InteractionType (split multi-values)
df["InteractionType"] = df["InteractionType"].str.split(",")
df = df.explode("InteractionType")
df["InteractionType"] = df["InteractionType"].str.strip()

```

```

df.columns = df.columns.str.strip()
df = df.drop([col for col in df.columns if "LoginBin_equal" in col], axis=1)

```

```

for col in ["Gender", "MaritalStatus", "IncomeLevel",
            "ServiceUsage", "InteractionType", "ResolutionStatus", "LoginBin_quartile"]:
    df[col] = df[col].astype(str).str.strip().str.lower()

# Replace string "nan" with actual NaN
df.replace("nan", np.nan, inplace=True)

# Check cleaned categories
for col in ["Gender", "MaritalStatus", "IncomeLevel", "ServiceUsage"]:
    print(col, df[col].unique())

```

```

Gender ['m' 'f']
MaritalStatus ['single' 'married' 'widowed' 'divorced']
IncomeLevel ['low' 'high' 'medium']
ServiceUsage ['mobile app' 'online banking' 'website']

```

```
df.head()
```

...	↑↓	C...	...	↑↓	...	↑↓	Marita...	...	↑↓	Inc...	...	↑↓	...	↑↓	Am...	...	↑↓	LoginFre...	...	↑↓	HasInter...	...	↑↓	ProductC...
0		1	m		single			low			62			416.5			34			1	Electronics			
1		10	m		married			high			68			1397.36			203			0	Furniture			
2		100	f		married			low			41			2217.12			360			1	Clothing			
3		100	f		married			low			41			2217.12			360			1	Clothing			
4		1000	m		widowed			low			34			1670.79			132			0	Furniture			

Rows: 5

 Expand

```
# Display unique values for all categorical (object) columns in df
for col in df.select_dtypes(include="object").columns:
    print(f"{col}: {df[col].unique()}\n")

Gender: ['m' 'f']

MaritalStatus: ['single' 'married' 'widowed' 'divorced']

IncomeLevel: ['low' 'high' 'medium']

ProductCategory: ['Electronics' 'Furniture' 'Clothing' 'Groceries' 'Books']

ServiceUsage: ['mobile app' 'online banking' 'website']

InteractionType: ['inquiry' 'no interaction' 'complaint' 'feedback']

ResolutionStatus: ['resolved' 'no interaction' 'unresolved']

ChurnStatusLabel: ['Active' 'Churned']

LoginBin_quartile: ['q1 (lowest)' 'q3' 'q4 (highest)' 'q2']
```

```
df.select_dtypes(exclude="object").columns

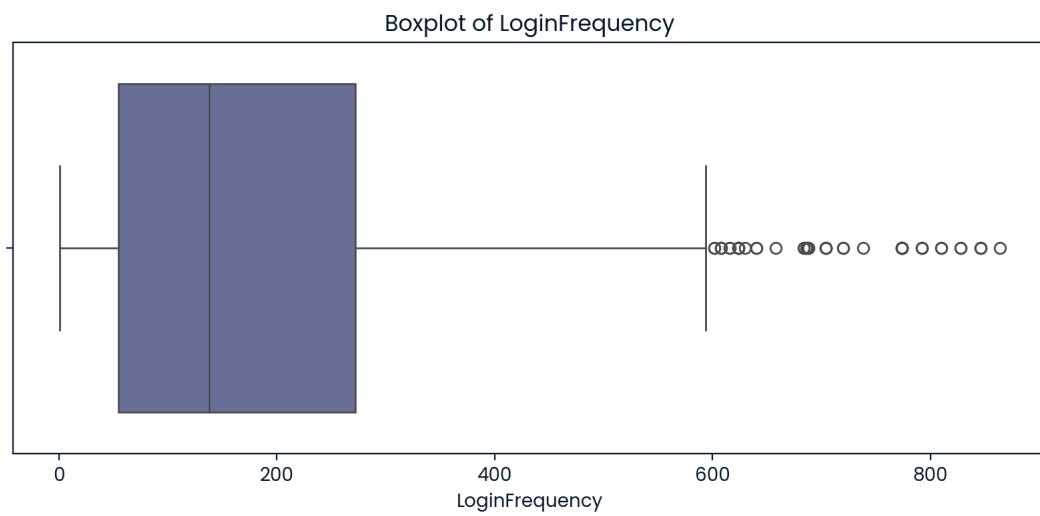
Index(['CustomerID', 'Age', 'AmountSpent', 'LoginFrequency', 'HasInteraction',
       'ChurnStatus'],
      dtype='object')
```

```
import matplotlib.pyplot as plt
import seaborn as sns

# Visualize LoginFrequency to spot outliers
plt.figure(figsize=(10, 4))
sns.boxplot(x=df["LoginFrequency"])
plt.title("Boxplot of LoginFrequency")
plt.show()

# Calculate outlier thresholds using IQR
Q1 = df["LoginFrequency"].quantile(0.25)
Q3 = df["LoginFrequency"].quantile(0.75)
IQR = Q3 - Q1
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

# Find outliers
outliers = df[(df["LoginFrequency"] < lower_bound) | (df["LoginFrequency"] > upper_bound)]
```



```
outliers = df[(df["LoginFrequency"] < lower_bound) | (df["LoginFrequency"] > upper_bound)]
outliers[["CustomerID", "LoginFrequency", "LoginBin_quartile"]].describe(include="all")
```

...	Custom...	...	LoginFre...	...	LoginBin_qua...	...	
count		48		48	48		
unique					1		
top					q4 (highest)		
freq					48		
mean	567.1458333333		712.8333333333		null		
std	312.9555735981		81.0545740328		null		
min		15		602	null		
25%		399		628.5	null		
50%		584		704	null		
75%		887		774	null		
max		986		864	null		

Rows: 11

Expand

Count: 48 customers flagged as outliers LoginFrequency range:

- Min: 602
- Max: 864
- Mean: 713
- Median: 704 These are not data errors — they're real customers with extremely high login activity. They represent a distinct segment: highly engaged users who may be loyal or at risk if their behavior changes. Their presence distorts the scale of LoginFrequency, but they're valuable for churn modeling.

```
print(df.columns.tolist())
```

```
['CustomerID', 'Gender', 'MaritalStatus', 'IncomeLevel', 'Age', 'AmountSpent', 'LoginFrequency', 'HasInteraction',
'ProductCategory', 'ServiceUsage', 'InteractionType', 'ResolutionStatus', 'ChurnStatus', 'ChurnStatusLabel',
'LoginBin_quartile']
```

```
df["flag_login_outlier"] = (df["LoginFrequency"] > 600).astype(int)
```

Write Python code or [tell our AI what to do](#)

```
print(df.columns.tolist())
```

```
['CustomerID', 'Gender', 'MaritalStatus', 'IncomeLevel', 'Age', 'AmountSpent', 'LoginFrequency', 'HasInteraction',
'ProductCategory', 'ServiceUsage', 'InteractionType', 'ResolutionStatus', 'ChurnStatus', 'ChurnStatusLabel',
'LoginBin_quartile', 'flag_login_outlier']
```

## Predictive Modeling for Customer Churn

In this notebook, we will develop and implement a predictive model using Random Forest and other machine learning algorithms to predict customer churn. Our goal is to achieve an ROC-AUC score of at least 0.82. We will use the `customer_level.csv` dataset and perform exploratory data analysis (EDA) as needed.

```
import pandas as pd
```

```
# Load the dataset
df = pd.read_csv('customer_level.csv')
df.head()
```

...	C...	...	...	...	Marita...	...	Inc...	...	...	Am...	...	...	...	...	...	...	...	...	...	...
0		1	M		Single		Low		62		416.5		34		1	Electronics				
1		10	M		Married		High		68		1397.36		203		0	Furniture				
2		100	F		Married		Low		41		2217.12		360		1	Clothing				
3		100	F		Married		Low		41		2217.12		360		1	Clothing				
4		1000	M		Widowed		Low		34		1670.79		132		0	Furniture				

Rows: 5

Expand

## Exploratory Data Analysis (EDA)

Let's explore the dataset to understand its structure, check for missing values, and get basic statistics.

```
# Check dataset info and missing values
df.info()
df.isnull().sum()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1215 entries, 0 to 1214
Data columns (total 16 columns):
 #   Column           Non-Null Count  Dtype  
---  --  
0   CustomerID      1215 non-null    int64  
1   Gender          1215 non-null    object  
2   MaritalStatus   1215 non-null    object  
3   IncomeLevel     1215 non-null    object  
4   Age              1215 non-null    int64  
5   AmountSpent     1215 non-null    float64 
6   LoginFrequency  1215 non-null    int64  
7   HasInteraction  1215 non-null    object  
8   ProductCategory 1215 non-null    object  
9   ServiceUsage    1215 non-null    object  
10  InteractionType 1215 non-null    object  
11  ResolutionStatus 1215 non-null    object  
12  ChurnStatus     1215 non-null    int64  
13  ChurnStatusLabel 1215 non-null    object  
14  LoginBin_equal  282 non-null    object  
15  LoginBin_quartile 1215 non-null    object  
dtypes: float64(1), int64(5), object(10)
memory usage: 152.0+ KB
```

index	...	↑↓	...	↑↓
CustomerID		0		
Gender		0		
MaritalStatus		0		
IncomeLevel		0		
Age		0		
AmountSpent		0		
LoginFrequency		0		
HasInteraction		0		
ProductCategory		0		
ServiceUsage		0		
InteractionType		0		
ResolutionStatus		0		
ChurnStatus		0		
ChurnStatusLabel		0		
LoginBin_equal		933		
LoginBin_quartile		0		

Rows: 16

Expand

# Display basic statistics df.describe(include='all')																						
...	↑↓	Custom...	...	↑↓	...	↑↓	Marita...	...	↑↓	Inc...	...	↑↓	Age	...	↑↓	AmountSp...	...	↑↓	LoginFre...	...	↑↓	HasInter...
count		1215	1215		1215					1215			1215			1215	1215		1215			
unique			2		4					3												
top			F		Widowed					High												
freq			630		322					423												
mean		496.8576131687	null		null					null			43.4082304527			1877.384781893			189.5094650206		0.726	
std		287.6945747523	null		null					null			15.2065817953			1346.8658953484			176.2837817209		0.445	
min		1	null		null					null			18			13.86			1			
25%		247.5	null		null					null			30			865.485			54.5			
50%		492	null		null					null			44			1562.83			138			
75%		748.5	null		null					null			56			2571.24			272			
max		1000	null		null					null			69			6440.6			864			

Rows: 11

Expand

## Data Preprocessing

We will preprocess the data by handling missing values, encoding categorical variables, and splitting the data into features and target.

```
# Example preprocessing (update as needed based on EDA results)
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder

# Assume 'Churn' is the target variable (update if different)
target_col = 'ChurnStatus'

# Encode categorical variables
df_encoded = df.copy()
for col in df_encoded.select_dtypes(include=['object', 'category']).columns:
    if col != target_col:
        df_encoded[col] = LabelEncoder().fit_transform(df_encoded[col].astype(str))

# Handle missing values (simple fill for demonstration)
df_encoded = df_encoded.fillna(df_encoded.median(numeric_only=True))

# Split features and target
X = df_encoded.drop([target_col, "ChurnStatusLabel"], axis=1)

y = df_encoded[target_col]

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42, stratify=y)
```

## Model Training and Evaluation

We will train a Random Forest classifier and compare its performance with other algorithms such as Logistic Regression and XGBoost. The main evaluation metric will be ROC-AUC.

```

from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from xgboost import XGBClassifier
from sklearn.metrics import roc_auc_score, roc_curve

# Train Random Forest
rf = RandomForestClassifier(random_state=42)
rf.fit(X_train, y_train)
rf_probs = rf.predict_proba(X_test)[:,1]
rf_auc = roc_auc_score(y_test, rf_probs)

# Train Logistic Regression
lr = LogisticRegression(max_iter=1000, random_state=42)
lr.fit(X_train, y_train)
lr_probs = lr.predict_proba(X_test)[:,1]
lr_auc = roc_auc_score(y_test, lr_probs)

# Train XGBoost
xgb = XGBClassifier(use_label_encoder=False, eval_metric='logloss', random_state=42)
xgb.fit(X_train, y_train)
xgb_probs = xgb.predict_proba(X_test)[:,1]
xgb_auc = roc_auc_score(y_test, xgb_probs)

print(f"Random Forest ROC-AUC: {rf_auc:.3f}")
print(f"Logistic Regression ROC-AUC: {lr_auc:.3f}")
print(f"XGBoost ROC-AUC: {xgb_auc:.3f}")

```

Random Forest ROC-AUC: 0.690  
 Logistic Regression ROC-AUC: 0.524  
 XGBoost ROC-AUC: 0.684

```

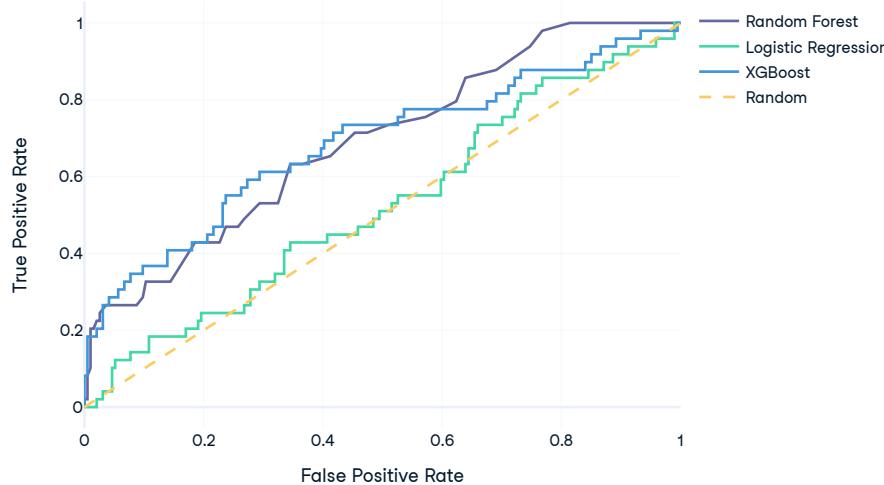
import plotly.graph_objs as go

# Plot ROC curves
fpr_rf, tpr_rf, _ = roc_curve(y_test, rf_probs)
fpr_lr, tpr_lr, _ = roc_curve(y_test, lr_probs)
fpr_xgb, tpr_xgb, _ = roc_curve(y_test, xgb_probs)

fig = go.Figure()
fig.add_trace(go.Scatter(x=fpr_rf, y=tpr_rf, mode='lines', name='Random Forest'))
fig.add_trace(go.Scatter(x=fpr_lr, y=tpr_lr, mode='lines', name='Logistic Regression'))
fig.add_trace(go.Scatter(x=fpr_xgb, y=tpr_xgb, mode='lines', name='XGBoost'))
fig.add_trace(go.Scatter(x=[0,1], y=[0,1], mode='lines', name='Random', line=dict(dash='dash')))
fig.update_layout(title='ROC Curves', xaxis_title='False Positive Rate', yaxis_title='True Positive Rate', width=700, height=500)
fig.show()

```

## ROC Curves



```
print(df_encoded.corr()[target_col].sort_values(ascending=False))
```

ChurnStatus	1.000000
ChurnStatusLabel	1.000000
HasInteraction	0.026155
Age	0.025463
IncomeLevel	0.019466
Gender	0.014467
ResolutionStatus	0.010732
ProductCategory	0.008132
MaritalStatus	0.001630
CustomerID	-0.011954
AmountSpent	-0.024454
InteractionType	-0.035769
LoginBin_equal	-0.065084
LoginBin_quartile	-0.065314
ServiceUsage	-0.068383
LoginFrequency	-0.072942
Name: ChurnStatus, dtype: float64	